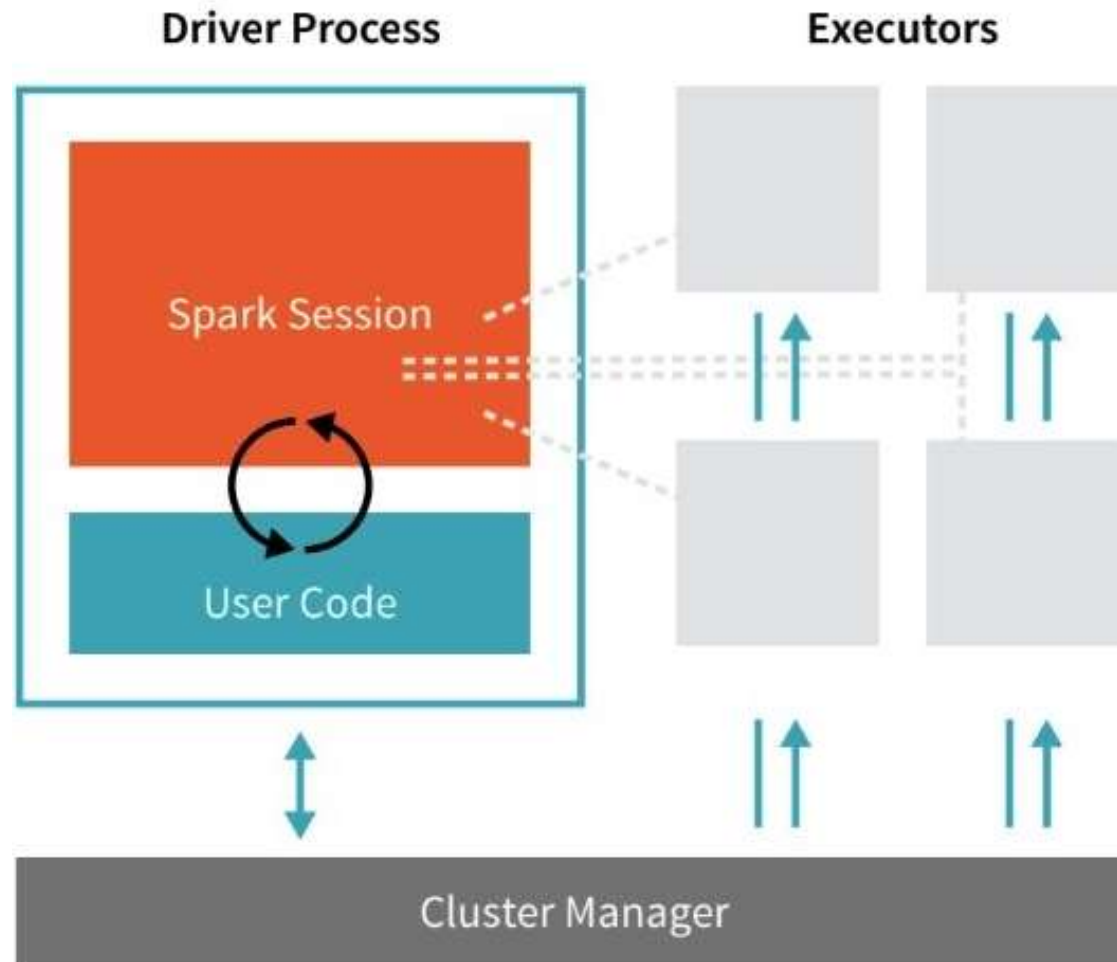



An Introduction to Apache Spark


Spark Architecture



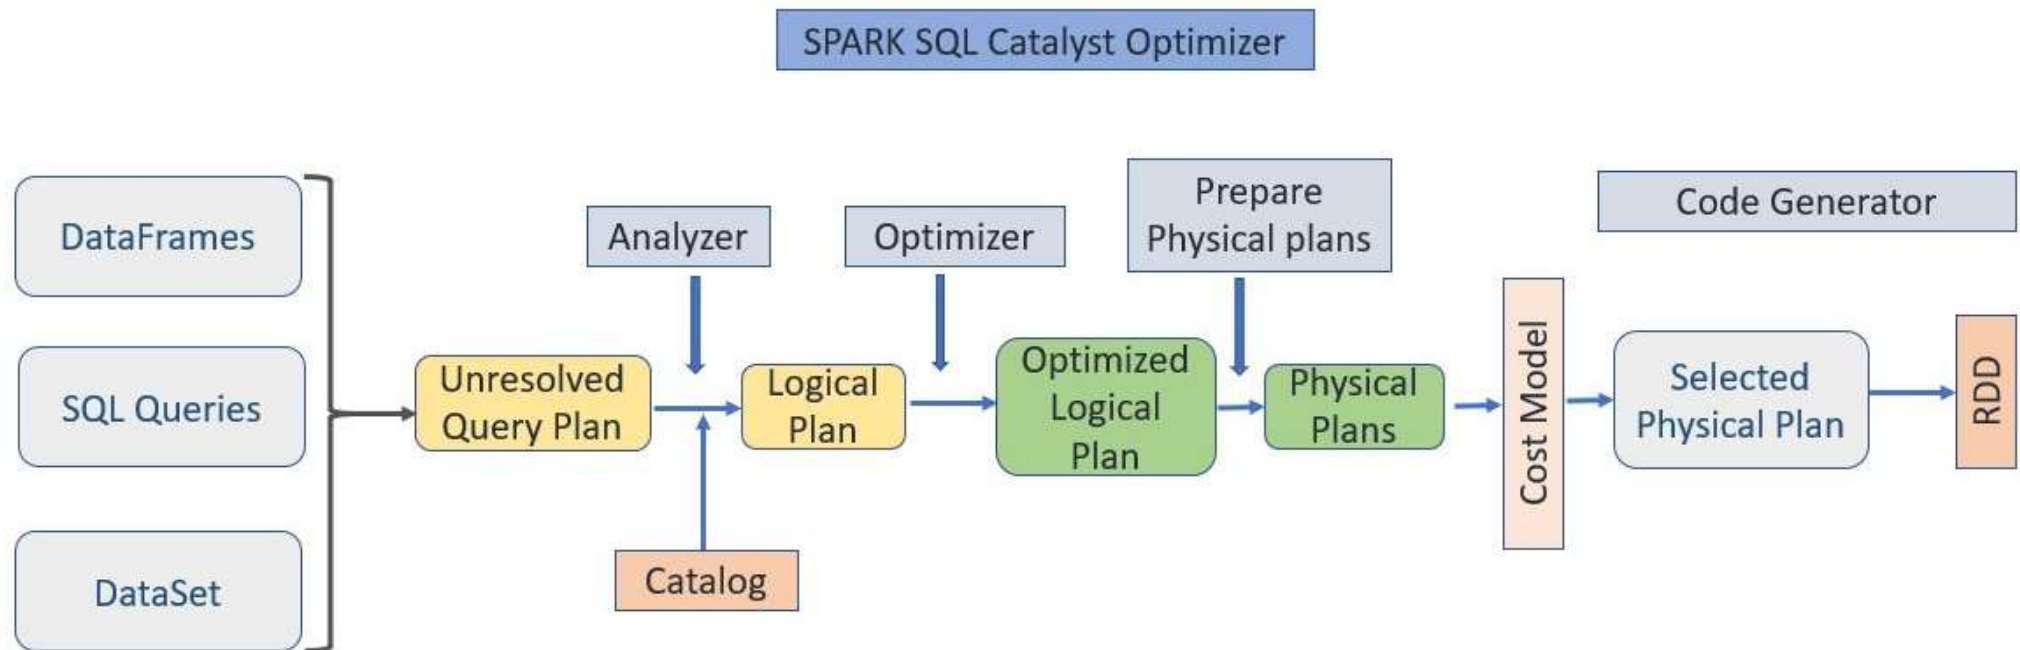
Spark's Language APIs



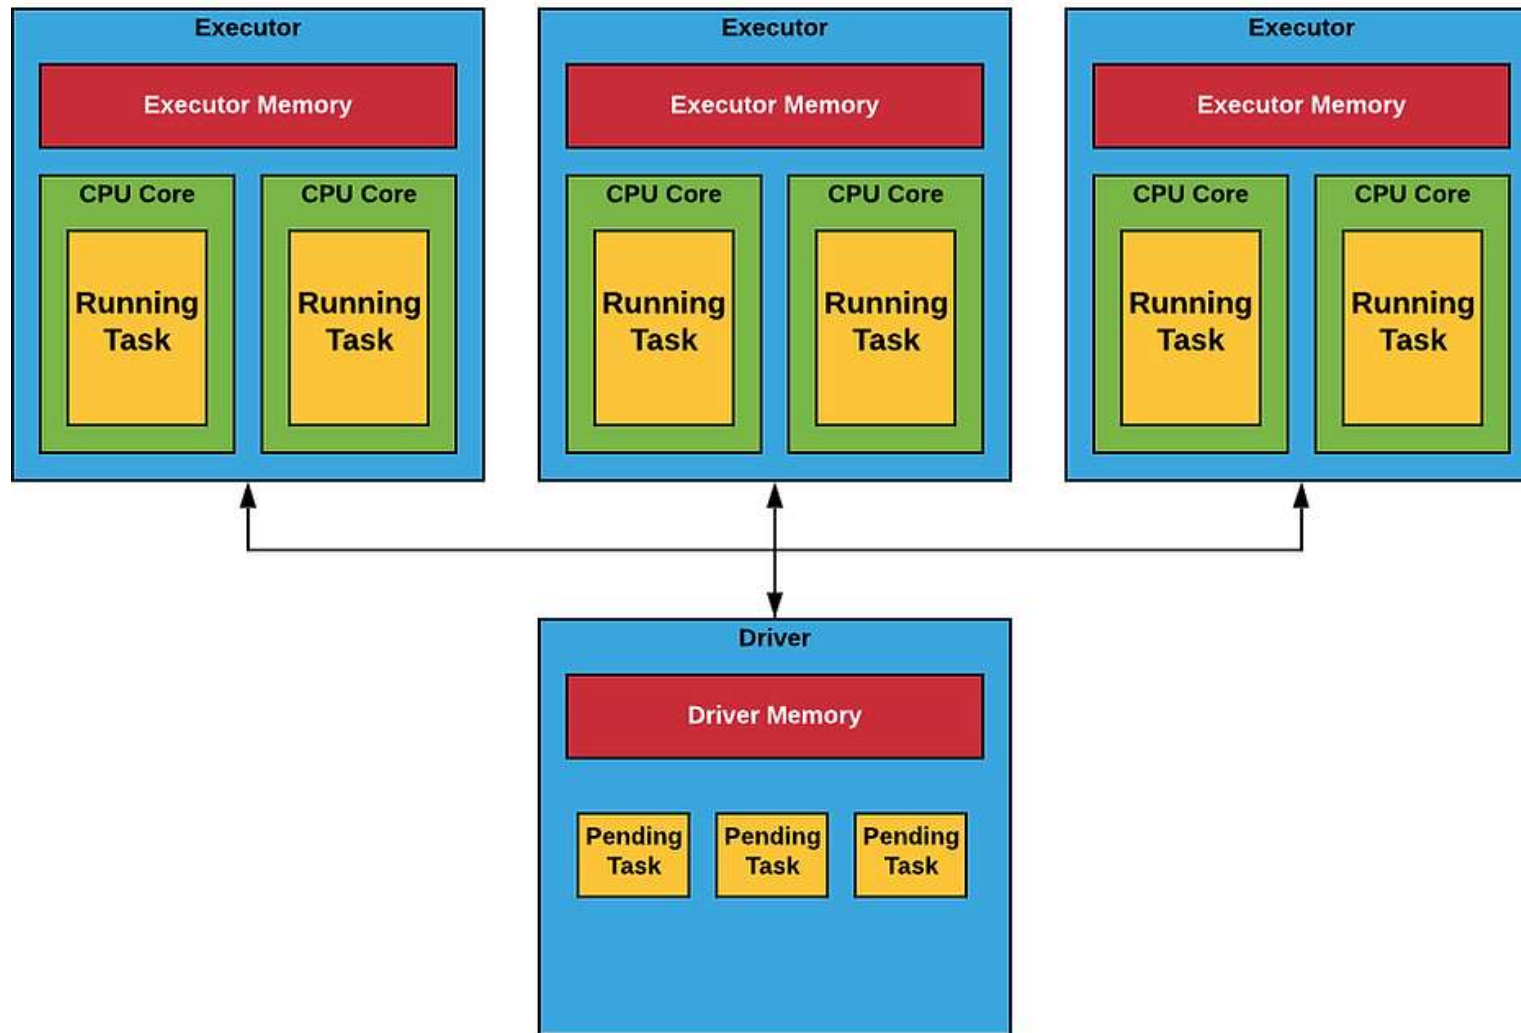
Ways to create DataFrame in Spark




Spark Optimization



Spark Executor / Cores / Memory



Partition and Executor



Repartitioning


- The data in the cluster will be split.
- Repartition involves a full shuffle
- Coalesce reduces the number of partitions and avoids a full shuffle

Partitioning on input stage

- When Spark reads a file from HDFS, it creates a single partition for a single input split

Data Size	Block Size	Total Blocks	Partitions
10240	128	80	80

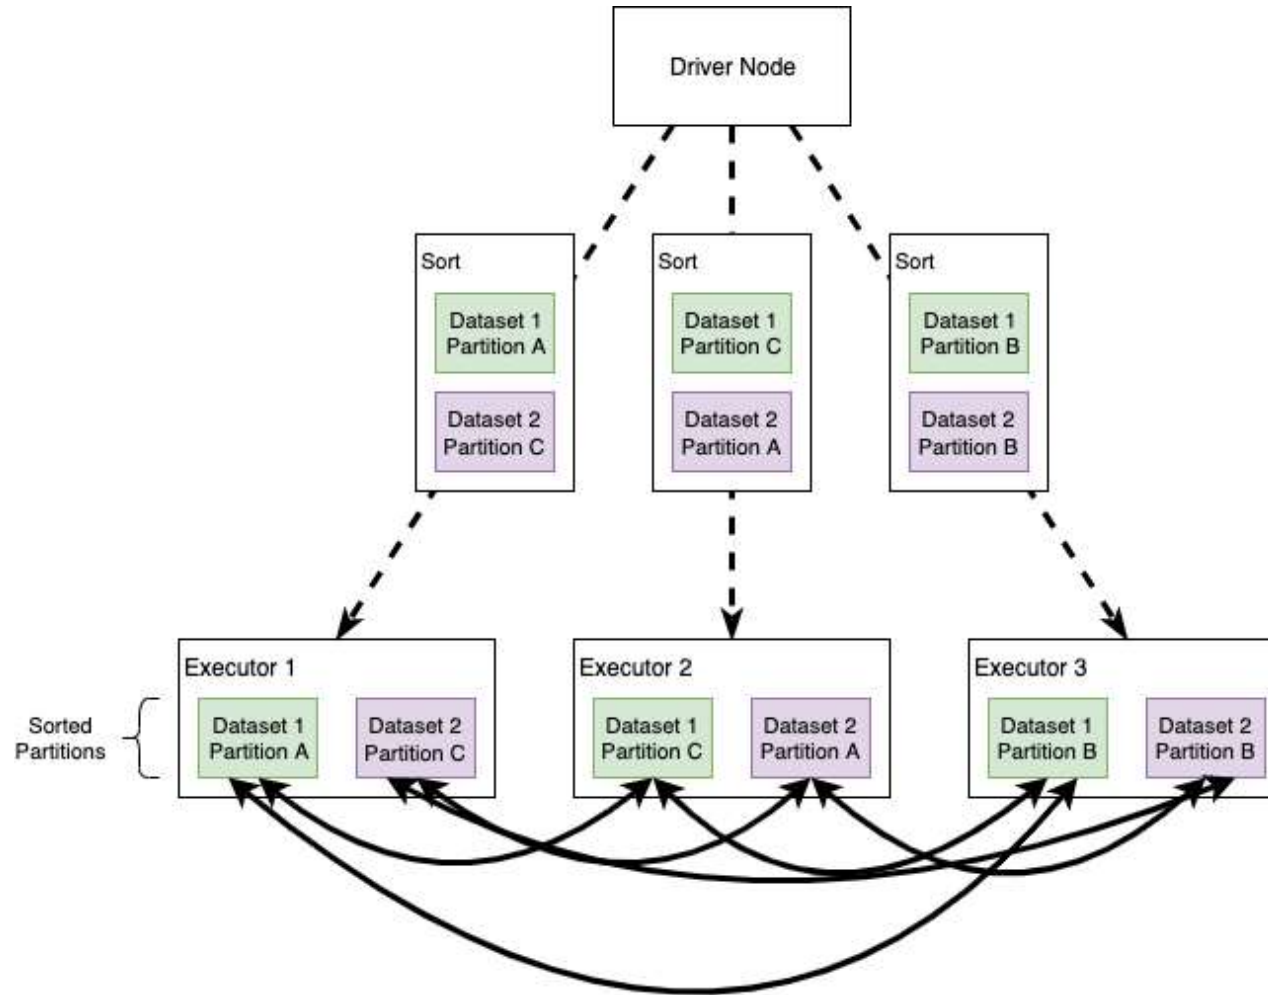
Shuffle partitioning



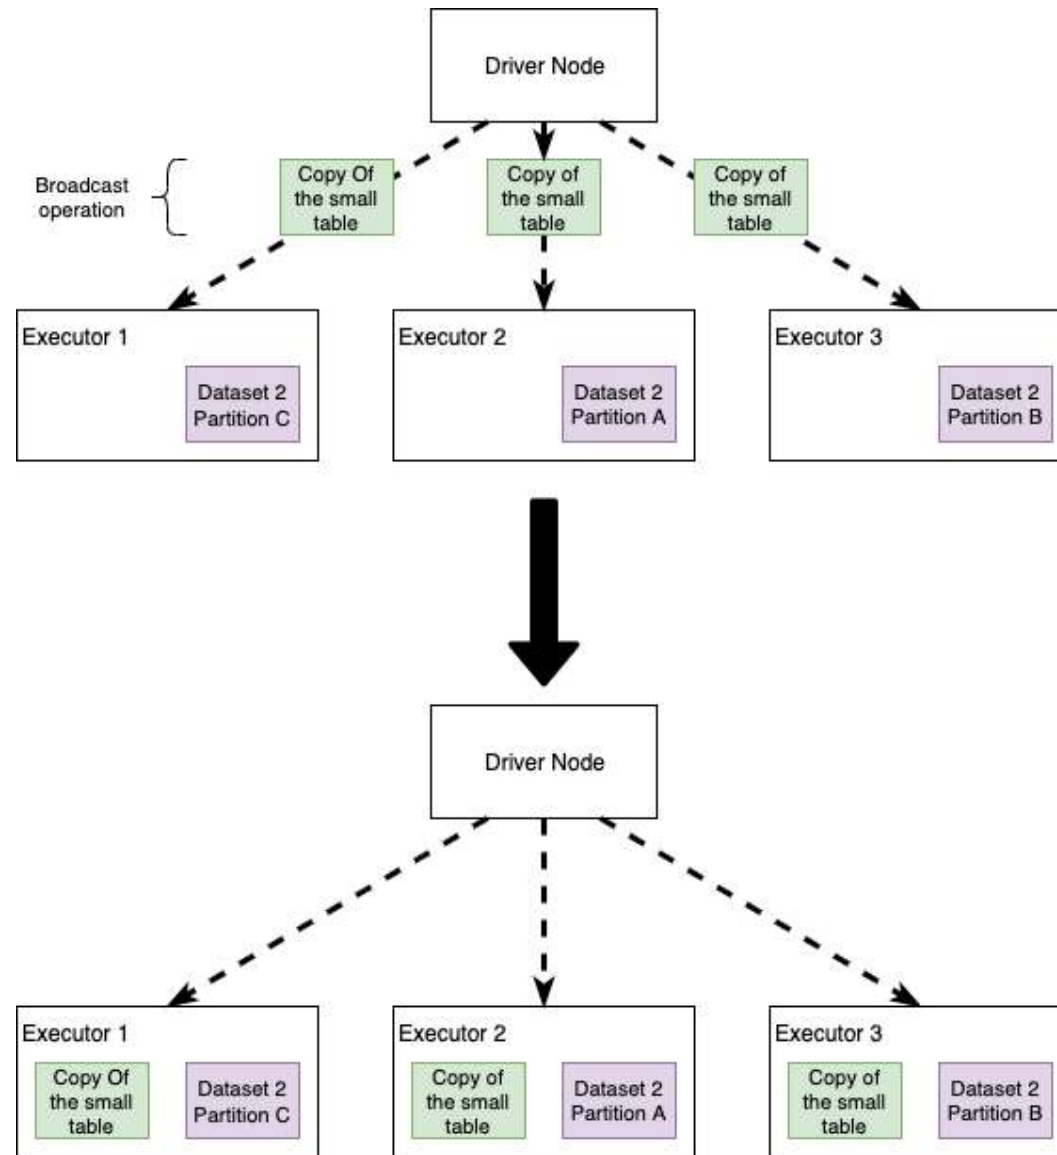
Output partitioning

- Saving the partitioned data on the properly selected condition can significantly speed up the reading and retrieval of data

Sort Merge Join



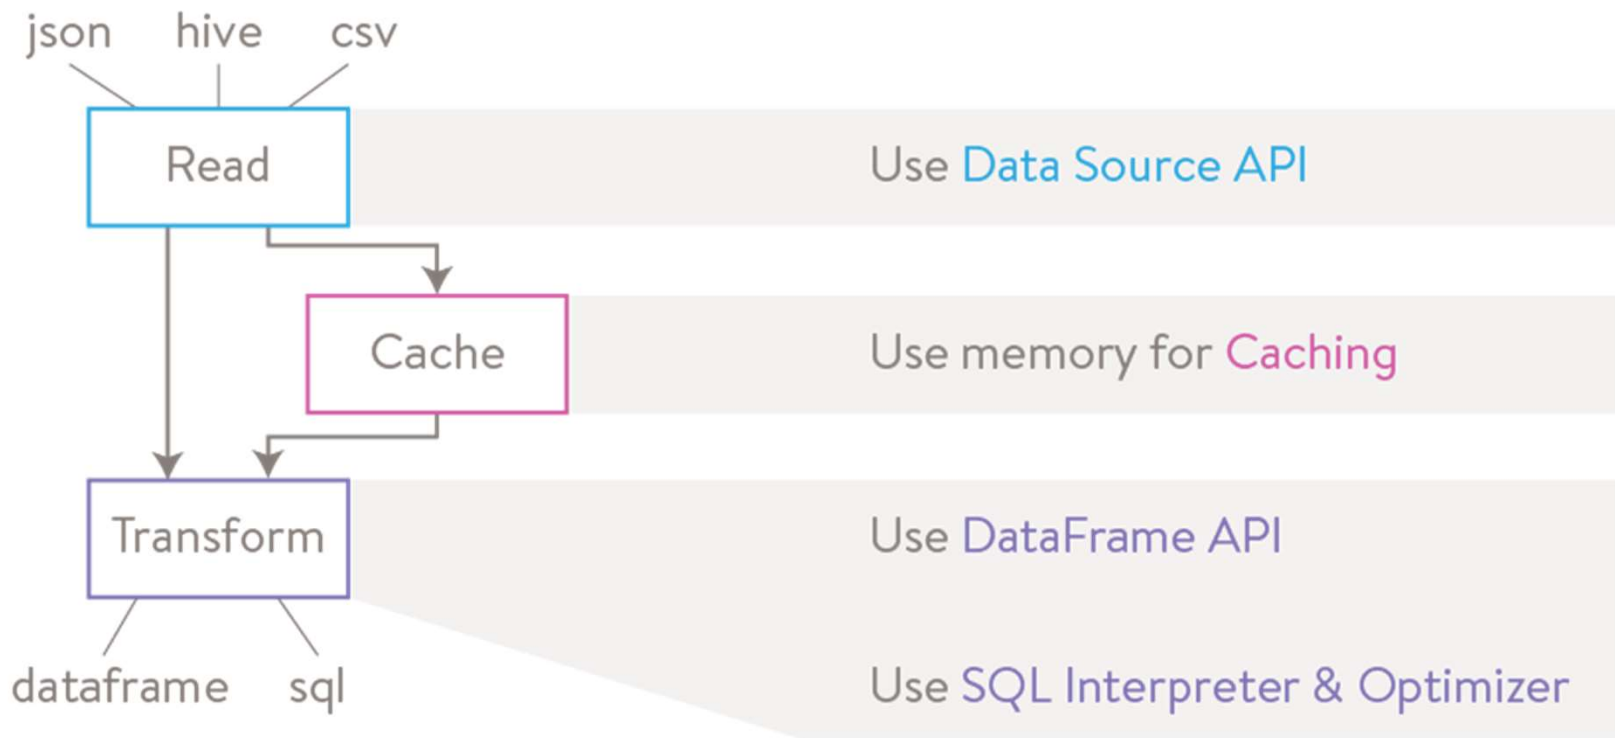
Broadcast Join



Caching

Data Processing Workflow

Strategies



Caching Storage Levels

1. MEMORY_ONLY (Default level)

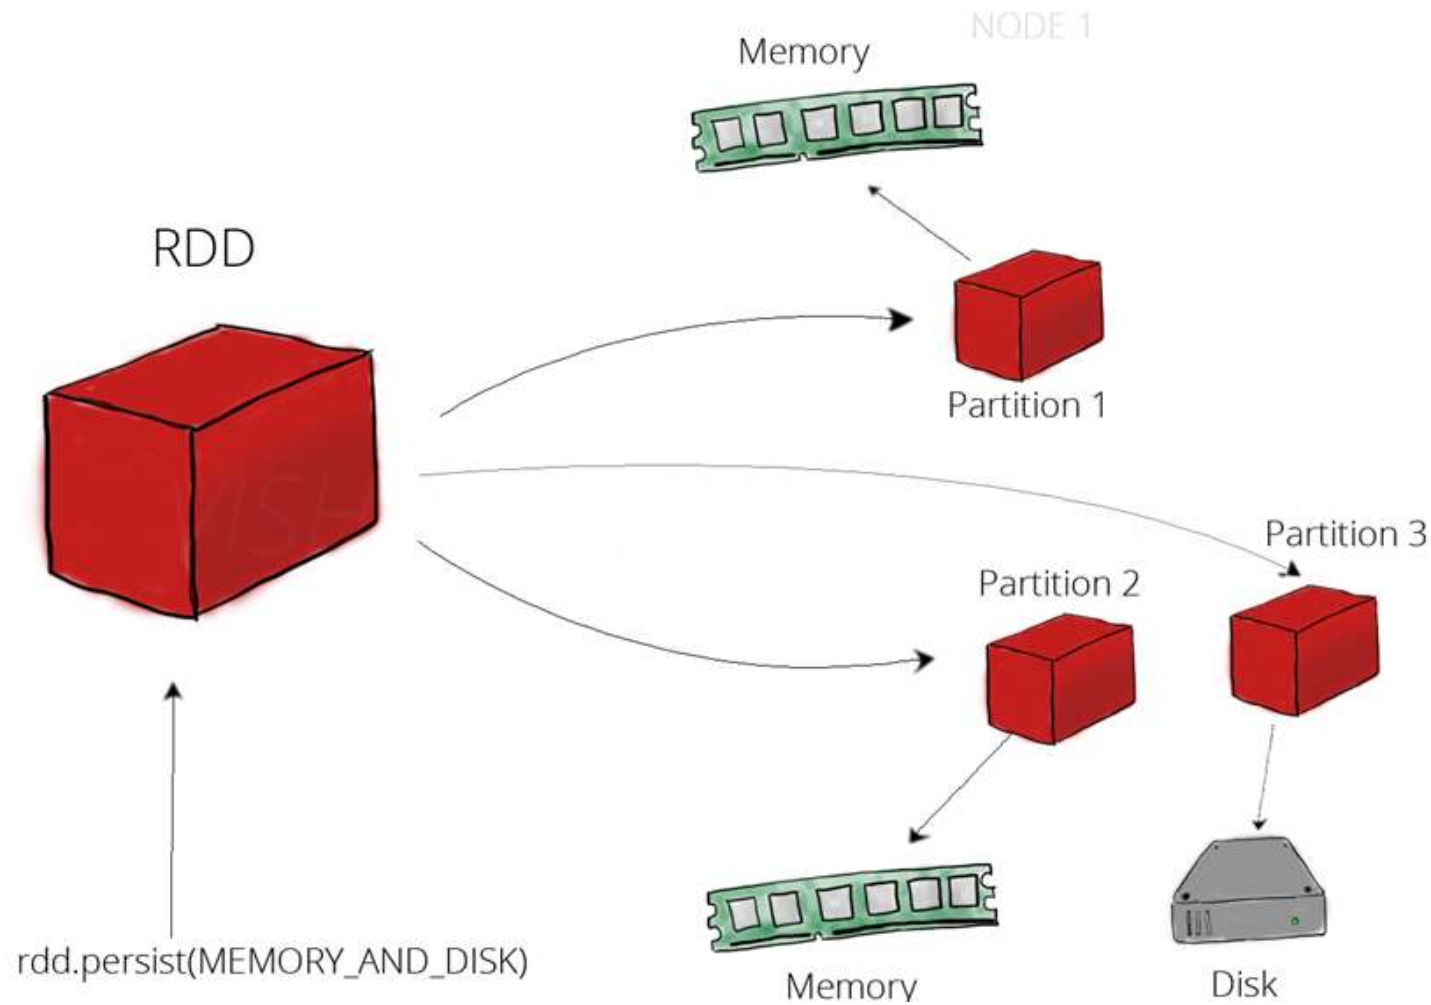
2. MEMORY_AND_DISK

3. MEMORY_ONLY_SER

4. MEMORY_ONLY_DISK_SER

5. DISC_ONLY

MEMORY_AND_DISK



MEMORY_ONLY_SER

- Store RDD as serialized Java objects
- More space-efficient

MEMORY_AND_DISK_SER

- Spill partitions that don't fit in memory to disk

Caching Storage Levels

Level	Space used	CPU time	In memory	On disk	Serialized

MEMORY_ONLY	High	Low	Y	N	N
MEMORY_ONLY_SER	Low	High	Y	N	Y
MEMORY_AND_DISK	High	Medium	Some	Some	Some
MEMORY_AND_DISK_SER	Low	High	Some	Some	Y
DISK_ONLY	Low	High	N	Y	Y

Unpersist RDD

- Spark drop out the old data partition in the LRU (least recently used) fashion.
 - LRU is an algorithm which ensures the least frequently used data
- Can also remove the cache manually using:
 - `.unpersist()`

Which Storage Level to Choose?

- If data fit comfortably with the default storage level (MEMORY_ONLY), leave them that way.
- If not, try using MEMORY_ONLY_SER
- Use the replicated storage levels if you want fast fault recovery

Thanks